

# **Cloud\lative** Lives

# K8s初体验-快速部署第一个容器应用

华为云容器团队倾心打造



# 大纲

- Kubernetes Pod详解
- Kubernetes工作负载与服务介绍
- 快速部署第一个容器应用







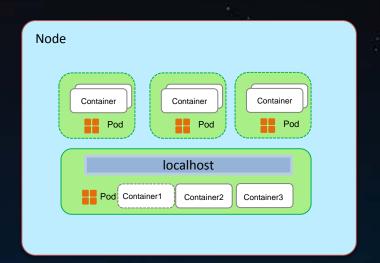


**HUAWEI** 

## Kubernetes关键概念-Pod







- 在Kubernetes中, pods是能够创建、调度、和管理的最小部署单元, 是一组容器的集合, 而 不是单独的应用容器
- 同一个Pod里的容器共享同一个网络命名空间,IP地址及端口空间。
- 从生命周期来说, Pod是短暂的而不是长久的应用。Pods被调度到节点, 保持在这个节点上 直到被销毁。

```
POD实例:
 "kind": "Pod",
 "apiVersion": "v1",
 "metadata": {
  "name": "redis-django",
  "labels": {
   "app": "webapp"
"spec": {
  "containers": [
    "name": "key-value-store",
    "image": "redis"
    "name": "frontend",
    "image": "django"
```





### Pod详解-容器 (Containers)



### 容器分类

- Infrastructure Container:基础容器
  - 用户不可见,无需感知
  - 维护整个Pod网络空间
- InitContainers:初始化容器,一般用于服务等待处理以及注册Pod信息等。
  - 先于业务容器开始执行
  - 顺序执行,执行成功退出(exit 0),全部执行成功后开始启动业务容器
- Containers: 业务容器
  - 并行启动,启动成功后一直Running

apiVersion: v1

kind: Pod metadata:

name: myapp-pod

labels:

app: myapp

spec:

#### containers:

- name: myapp-container

image: busybox

command: ['sh', '-c', 'echo The app is running!

&& sleep 3600']

#### initContainers:

- name: init-myservice

image: busybox

command: ['sh', '-c', 'until nslookup myservice;

do echo waiting for myservice; sleep 2; done;']

name: init-mydb image: busybox

command: ['sh', '-c', 'until nslookup mydb; do

echo waiting for mydb; sleep 2; done;']









#### 镜像部分:

- 镜像地址和拉取策略
- 拉取镜像的认证凭据

#### 启动命令:

- command:替换docker容器的 entrypoint
- args:作为docker容器entrypoint的 入参

#### 计算资源:

- 请求值:调度依据
- 限制值:容器最大能使用的规格

```
spec:
 imagePullSecrets:
 - name: default-secret
 containers:
 - image: kube-dns:1.0.0
  imagePullPolicy: IfNotPresent
  command:
  - /bin/sh
  - -C
  - /kube-dns 1>>/var/log/skydns.log 2>&1 --domain=cluster.local. --dns-
port=10053
   --config-dir=/kube-dns-config --v=2
  resources:
   limits:
    cpu: 100m
    memory: 512Mi
   requests:
    cpu: 100m
    memory: 100Mi
```







### Pod详解-健康检查

#### 健康检查:

分命令行方式、httpGet请求方式以及TCPSocket方式。

- 业务探针(readinessProbe)
  - 探测不正常后,不会重启容器,只会拿掉服务后端的 endpoints
- 存活探针(livenessProbe)
  - 探测不正常后,会重启容器

```
spec:
 containers:
 - livenessProbe:
   failureThreshold: 5
   exec:
    command:
    - "/bin/sh",
    - "-c"
    - "echo 'iam ok'"
   initialDelaySeconds: 60
   periodSeconds: 10
   successThreshold: 1
   timeoutSeconds: 5
  readinessProbe:
   failureThreshold: 3
   httpGet:
    path: /readiness
    port: 8081
    scheme: HTTP
   initialDelaySeconds: 3
   periodSeconds: 10
   successThreshold: 1
   timeoutSeconds: 5
```

### Pod详解-外部输入

Pod可以接收的外部输入方式:环境变量、配置文件以及密 钥。

环境变量:使用简单,但一旦变更后必须重启容器。

- Key-value自定义
- From 配置文件 (configmap)
- From 密钥 (Secret)

以卷形式挂载到容器内使用,权限可控。

- 配置文件 (configmap)
- 密钥(secret)



secretName: secret







## 配置文件(ConfigMap)及密钥(Secret)介绍



• 大小不能超过1M

• Secret有类型区分,不同类型有不同校验方式,输入的value需base64加密

apiVersion: v1 data:

user: wangbo

details: |

XXX

kind: ConfigMap

metadata:

name: config

apiVersion: v1

data:

username: Y3Bvc3RncmVz

password: cGNwcGNw

kind: Secret metadata:

name: secret type: Opaque







### Pod详解-持久化存储



主机hostpath

必须要求Pod调度到固定节点上



- 云存储
  - 多类型选择,如云硬盘,文件存储,对象存储等
  - 不用担心Pod迁移

#### spec:

#### containers:

- image: tomcat:9-slim

imagePullPolicy: IfNotPresent

name: container-0 volumeMounts:

- mountPath: /usr/local/disk

name: pvc-153085714479448980

- mountPath: /usr/local/host

name: host

#### volumes:

- name: pvc-153085714479448980

persistentVolumeClaim:

claimName: pvc-153085714479448980

- name: host hostPath: path: /opt/







### PV/PVC介绍

apiVersion: v1

kind: PersistentVolumeClaim

metadata:

annotations:

volume.beta.kubernetes.io/storage-class: sata volume.beta.kubernetes.io/storage-provisioner:

flexvolume-huawei.com/fuxivol

labels:

failure-domain.beta.kubernetes.io/region: southchina failure-domain.beta.kubernetes.io/zone: az1.dc1

name: pvc-153085714479448980

spec:

accessModes:

- ReadWriteMany

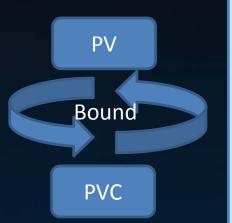
resources:

requests:

storage: 10Gi

volumeName: pv-b3377f21-80e2-11e8-ad55-

fa163e5a2cf6



apiVersion: v1

kind: PersistentVolume

metadata:

labels:

failure-domain.beta.kubernetes.io/region: southchina

failure-domain.beta.kubernetes.io/zone: az1.dc1 name: pv-b3377f21-80e2-11e8-ad55-fa163e5a2cf6

spec:

accessModes:

- ReadWriteMany

capacity:

storage: 10Gi

claimRef:

apiVersion: v1

kind: PersistentVolumeClaim

name: pvc-153085714479448980

flexVolume:

driver: huawei.com/fuxivol

fsType: ext4 options:

fsType: ext4

kubernetes.io/namespace: default

volumeID: 840d1d47-6303-4dd7-9a7f-41f4f80dbc1e

storageClassName: sata



创建 **PVC** 

设置到 Pod

通过PVC 找到PV

调度Pod到 合适的节点 根据PV挂载 存储到容器





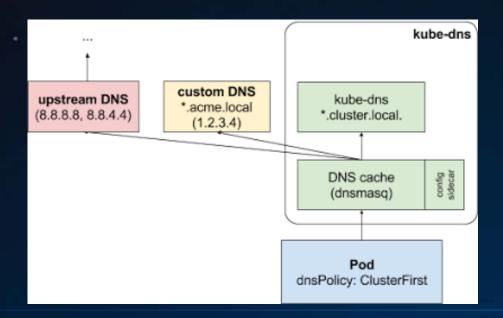




## Pod详解-服务域名发现



- dnsPolicy: Pod内域名解析的策略
  - ClusterFirst:使用kube-dns作为域名解析服务器
  - Default:使用节点(kubelet)指定的域名服务器解析域名
     器解析域名
  - ClusterFirstWithHostNet: 当Pod使用主机网络部署时使用



#### spec:

#### containers:

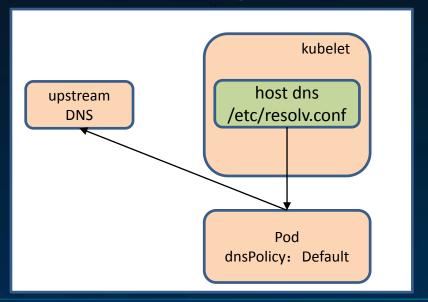
- command:
- /bin/sh
- -C
- stability --metrics=networkio#3000

image: stability:1.0

imagePullPolicy: IfNotPresent

name: stability

dnsPolicy: ClusterFirst









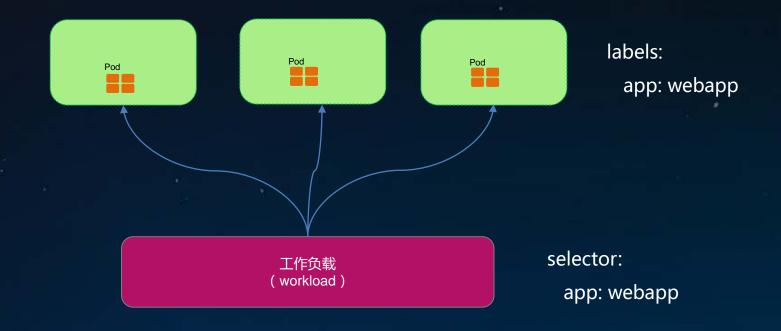
# 大纲

- Kubernetes Pod详解
- Kubernetes工作负载与服务介绍
- 快速部署第一个容器应用

## Pod与工作负载的关系

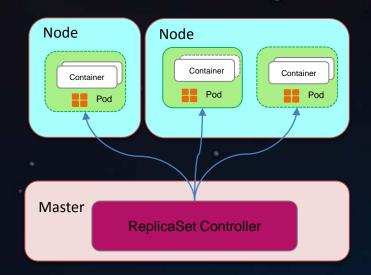


- 通过label-selector相关联
- Pod通过工作负载实现应用的运维,如伸缩、升级等





## 关键工作负载-ReplicaSet



- ReplicaSet—副本控制器
- 确保Pod的一定数量的份数(replica)在运行。如果超过这个数量,控制器会杀死一些,如果少了,控制器会启动一些。
- ReplicaSet用于解决pod的扩容和缩容问题。
- 通常用于无状态应用

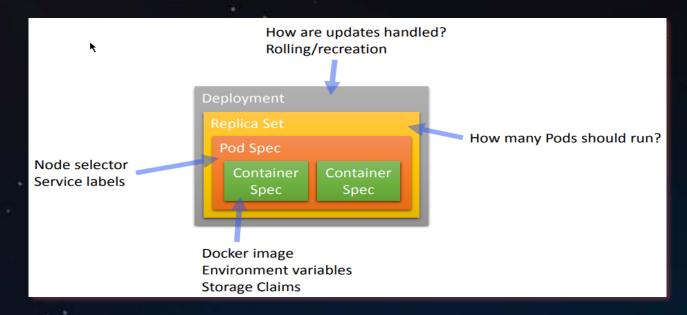
```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
 name: frontend
spec:
replicas: 3
selector:
  matchLabels:
   tier: frontend
  matchExpressions:
   - {key: tier, operator: In, values: [frontend]}
 template:
  metadata:
   labels:
    app: guestbook
    tier: frontend
  spec:
   containers:
   - name: php-redis
    image: gcr.io/google samples/gb-
frontend:v3
    resources:
     requests:
      cpu: 100m
      memory: 100Mi
    env:
    - name: GET HOSTS FROM
     value: dns
    ports:
    - containerPort: 80
```





HUAWEI

### 关键工作负载-Deployment



- Kubernetes Deployment提供了官方的用于更新Pod和Replica Set(下一代的Replication
   Controller)的方法,您可以在Deployment对象中只描述您所期望的理想状态(预期的运行状态),
   Deployment控制器为您将现在的实际状态转换成您期望的状态;
- Deployment集成了上线部署、滚动升级、创建副本、暂停上线任务,恢复上线任务,回滚到以前某一版本(成功/稳定)的Deployment等功能,在某种程度上,Deployment可以帮我们实现无人值守的上线,大大降低我们的上线过程的复杂沟通、操作风险。
- Deployment的典型用例:
  - 使用Deployment来启动 (上线/部署)一个Pod或者ReplicaSet
  - 检查一个Deployment是否成功执行
  - 更新Deployment来重新创建相应的Pods(例如,需要使用一个新的Image)
  - 如果现有的Deployment不稳定,那么回滚到一个早期的稳定的Deployment版本

apiVersion: extensions/v1beta1

kind: Deployment

metadata:

name: nginx-deployment

spec:

replicas: 3 template: metadata: labels:

app: nginx

spec:

containers:

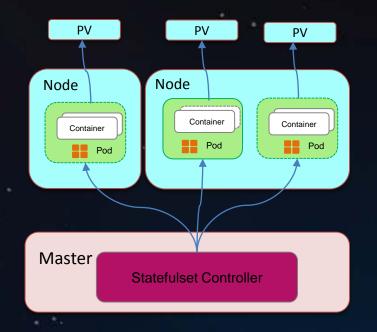
- name: nginx

image: nginx:1.7.9

ports:

- containerPort: 80

### 关键工作负载-StatefulSet



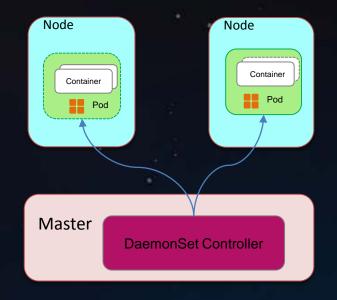
apiVersion: v1
kind: Service
metadata:
name: nginx
labels:
app: nginx
spec:
ports:
- port: 80
name: web
clusterIP: None
selector:
app: nginx

- StatefulSet—有状态应用
- 用于解决各个pod实例独立生命周期管理,提供各个实例的启动顺序和唯一性
  - 稳定,唯一的网络标识符。
  - 稳定,持久存储。
  - 有序的,优雅的部署和扩展。
  - 有序,优雅的删除和终止。
  - 有序的自动滚动更新。

```
kind: StatefulSet
metadata:
 name: web
spec:
 selector:
  matchLabels:
   app: nginx # has to
match .spec.template.metadata.labels
 serviceName: "nginx"
 replicas: 3 # by default is 1
 template:
  metadata:
   labels:
    app: nginx # has to
match .spec.selector.matchLabels
  spec:
   terminationGracePeriodSeconds: 10
   containers:
   - name: nginx
    image: k8s.gcr.io/nginx-slim:0.8
    ports:
    - containerPort: 80
     name: web
    volumeMounts:
    - name: www
     mountPath: /usr/share/nginx/html
 volumeClaimTemplates:
 - metadata:
   name: www
  spec:
   accessModes: [ "ReadWriteOnce" ]
   storageClassName: "mv-storage-class"
```

### 关键工作负载-DaemonSet





- DaemonSet能够让所有(或者一些特定)的Node节点运行同一个pod。当节点加入到kubernetes集 群中, pod会被(DaemonSet)调度到该节点上运行, 当节点从kubernetes集群中被移除,被 (DaemonSet)调度的pod会被移除,如果删除DaemonSet,所有跟这个DaemonSet相关的pods 都会被删除。
- 在使用kubernetes来运行应用时,很多时候我们需要在一个区域(zone)或者所有Node上运行同一 个守护进程(pod),例如如下场景:
  - 每个Node上运行一个分布式存储的守护进程,例如glusterd,ceph
  - 运行日志采集器在每个Node上,例如fluentd, logstash
  - 运行监控的采集端在每个Node,例如prometheus node exporter, collectd等

apiVersion: extensions/v1beta1

kind: DaemonSet

metadata:

name: tomcat-ds

spec:

template: metadata: labels:

app: tomcat

spec:

containers:

- name: tomcat

image: tomcat:8.0.30-jre8

ports:

- containerPort: 8080







## 关键工作负载-Job



分为普通任务 (Job) 和定时任务 (CronJob)

- 一次性执行,非堵塞
- 适合CI,视频解码,基因测序等业务

```
apiVersion: batch/v1
kind: Job
metadata:
name: pi
spec:
 template:
  metadata:
  name: pi
  spec:
  containers:
  - name: pi
    image: perl
    command: ["perl", "-Mbignum=bpi", "-
wle", "print bpi(2000)"]
   restartPolicy: Never
   backoffLimit: 4
```





### <u>Kubernetes关键概念</u>-CustomResourceDefinition



用户可自定义资源模型,可扩展行强:

• 需实现自己的controller进行自定义资源的管理

apiVersion: apiextensions.k8s.io/v1beta1 kind: CustomResourceDefinition metadata: # name must match the spec fields below, and be in the form: <plu>plural>.<group> name: crontabs.stable.example.com spec: # group name to use for REST API: /apis/<group>/<version> group: stable.example.com # list of versions supported by this CustomResourceDefinition versions: - name: v1 # Each version can be enabled/disabled by Served flag. served: true # One and only one version must be marked as the storage version. storage: true # either Namespaced or Cluster scope: Namespaced names: # plural name to be used in the URL: /apis/<group>/<version>/<plural> plural: crontabs # singular name to be used as an alias on the CLI and for display singular: crontab # kind is normally the CamelCased singular type. Your resource manifests use this. kind: CronTab # shortNames allow shorter string to match your resource on the CLI shortNames: - ct







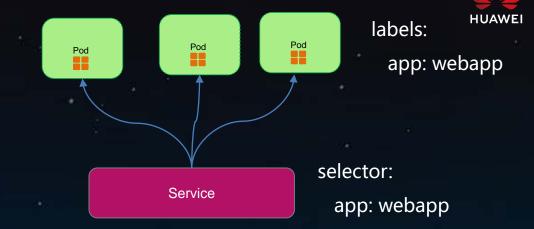
## Pod与服务的关系

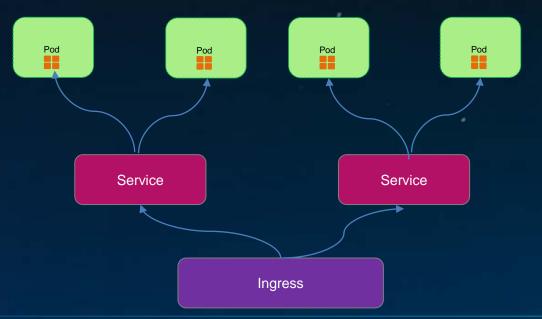
通过label-selector相关联

• 通过服务实现Pod的负载均衡能力:

Service:实现TCP/UDP 4层负载均衡能力

Ingress:实现HTTP 7层负载均衡能力

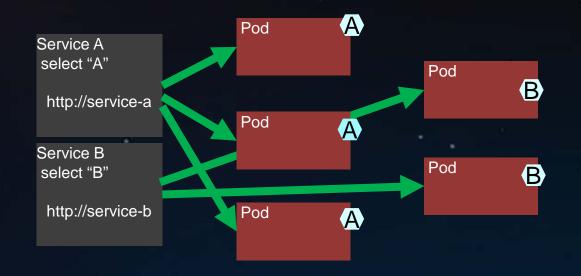






### Service





- Service定义了pods的逻辑集合和访问这个集合的策略。Pods集合是通过定义Service时提供的Label选择器完成的
- Service的引入旨在保证pod的动态变化对访问端透明,访问端只需要知道service的地址,由service来提供代理
- Service的抽象使得前端客户和后端Pods进行了解耦
- 支持ClusterIP, NodePort以及LoadBalancer三种类型
- Service的底层实现有userspace、iptables和ipvs三种模式

kind: Service apiVersion: v1 metadata:

name: my-service

spec:

selector: app: A ports:

protocol: TCP

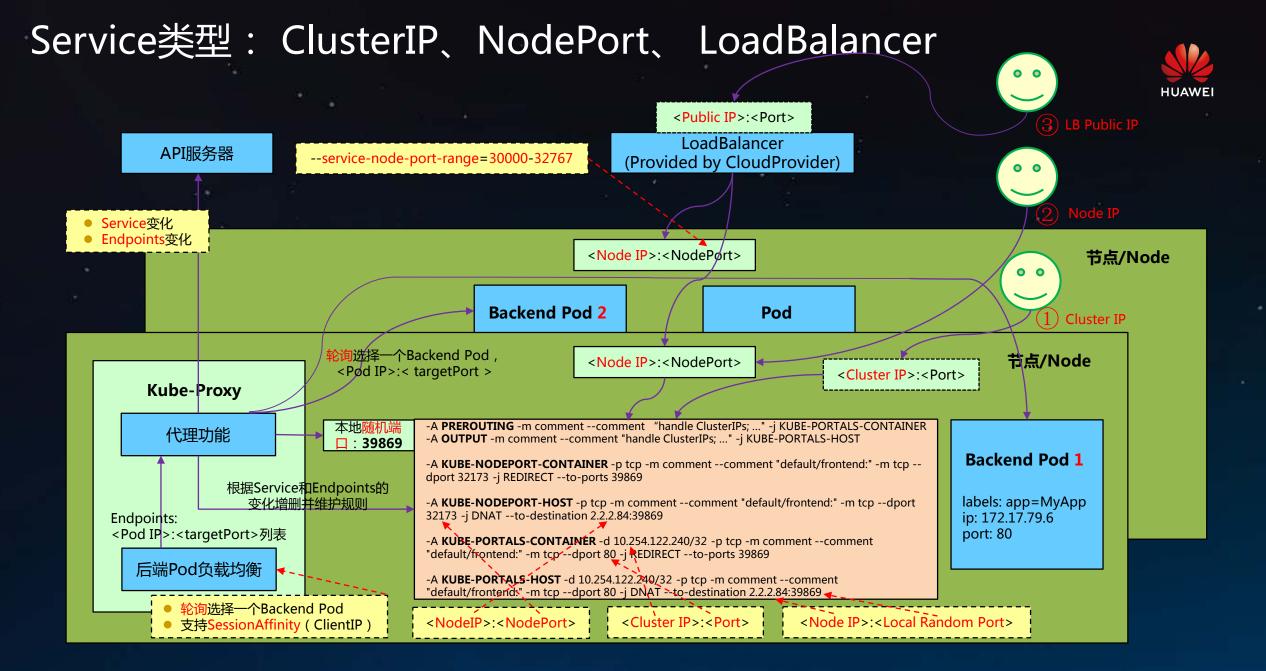
port: 80

targetPort: 9376 nodePort: 30061 type: ClusterIP







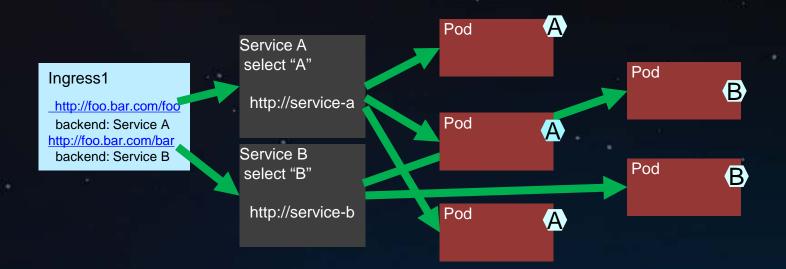






### Ingress





Ingress基于service实现7层路由转发能力

apiVersion: extensions/v1beta1 kind: Ingress metadata: name: test annotations: ingress.kubernetes.io/rewrite-target:/ spec: rules: - host: foo.bar.com http: paths: - path: /foo backend: serviceName: service-a servicePort: 80 - path: /bar backend: serviceName: service-b

servicePort: 80



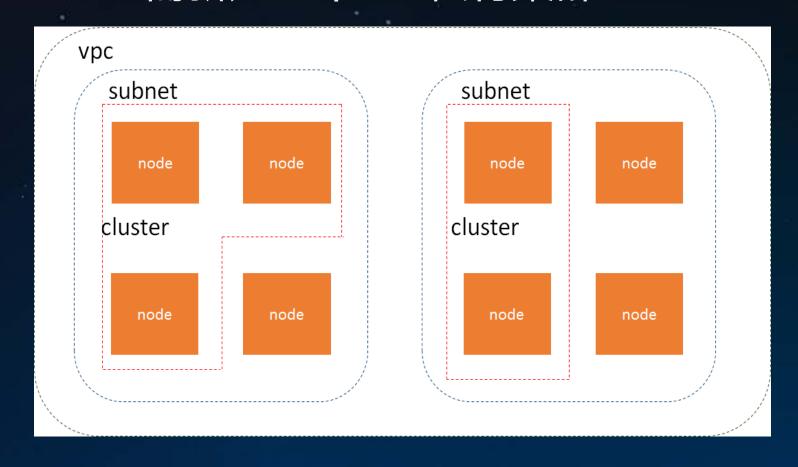
# 大纲

- Kubernetes Pod详解
- Kubernetes工作负载与服务介绍
- 快速部署第一个容器应用

## 在云上快速部署一个容器应用



## 需新建一个k8s私有集群





# Thank You

http://zhibo.huaweicloud.com/watch/2174406

直播 每周四 晚20:00





