
K8s 技术交流群

和宽哥一对一交流



课程推荐，提供跪舔式售后服务：

K8s 高薪全栈架构师课程： <https://edu.51cto.com/sd/518e5>

K8s CKA 认证课程： <https://edu.51cto.com/sd/fbbc8>

K8s CKS 认证课程： <https://edu.51cto.com/sd/affb5>

K8s 全栈架构师+CKA 套餐： <https://edu.51cto.com/topic/4973.html>

超级套购：

K8s 全栈架构师+CKA+CKS： <https://edu.51cto.com/topic/5174.html>

关注宽哥：

<https://edu.51cto.com/lecturer/11062970.html?type=2>

考题由简单到难排列

Kubectl 命令大全 <https://kubernetes.io/zh/docs/reference/kubectl/cheatsheet/> <https://edu.51cto.com/lecturer/11062970.html?type=2>

考试时先挑简单的做~~~

1.1 监控 Pod 日志

[< Previous](#)[☰ Task 14 of 17 ▾](#)[→ Next](#)

Task weight: 5%

Set configuration context:



```
[student@node-1] $ | kubectl config use-context  
k8s
```

Task

Monitor the logs of pod `foobar` and:

- Extract log lines corresponding to error `unable-to-access-website`
- Write them to `/opt/KUTR00101/foobar`

中文解释:

监控名为 `foobar` 的 Pod 的日志，并过滤出具有 `unable-access-website` 信息的行，然后将写入到 `/opt/KUTR00101/foobar`

解题:

```
$ kubectl config use-context k8s  
$ kubectl logs foobar | grep unable-access-website >  
/opt/KUTR00101/foobar
```

1.2 监控 Pod 度量指标

Task weight: 5%

Set configuration context:



```
[student@node-1] $ | kubectl config use-context  
k8s
```

Task

From the pod label `name=cpu-user`, find pods running high CPU workloads and write the name of the pod consuming most CPU to the file `/opt/KUTR00401/KUTR00401.txt` (which already exists).

中文解释：

找出具有 `name=cpu-user` 的 Pod，并过滤出使用 CPU 最高的 Pod，然后把它的名字写在已经存在的 `/opt/KUTR00401/KUTR00401.txt` 文件里（注意他没有说指定 `namespace`。所以需要使用 `-A` 指定所以 `namespace`）

解题：

```
$ kubectl config use-context k8s  
$ kubectl top po -A -l name=cpu-user  
NAMESPACE      NAME                                CPU (cores)    MEMORY (bytes)  
kube-system     coredns-54d67798b7-hl8xc           7m             8Mi  
kube-system     coredns-54d67798b7-m4m2q           6m             8Mi  
# 注意这里的 pod 名字以实际名字为准，按照 CPU 那一列进行选择一个最大的 Pod，另外如果  
CPU 的数值是 1 2 3 这样的。是大于带 m 这样的，因为 1 颗 CPU 等于 1000m，注意要用 >> 而不是 >
```

```
$ echo "coredns-54d67798b7-hl8xc" >> /opt/KUTR00401/KUTR00401.txt
```

1.3 Deployment 扩缩容

[< Previous](#)[☰ Task 8 of 17 ▾](#)[→ Next](#)

Task weight: 4%

Set configuration context:



```
[student@node-1] $ | kubectl config use-context  
k8s
```

Task

Scale the deployment **loadbalancer** to **6** pods.

中文解释:

扩容名字为 loadbalancer 的 deployment 的副本数为 6

解题:

```
$ kubectl config use-context k8s
```

```
$ kubectl scale --replicas=6 deployment loadbalancer
```

```
$ kubectl edit
```

1.4 检查 Node 节点的健康状态

Set configuration context:



```
[student@node-1] $ | kubectl config use-context  
k8s
```

Task

Check to see how many nodes are ready (not including nodes tainted `NoSchedule`) and write the number to `/opt/KUSC00402/kusc00402.txt` .

中文解释：

检查集群中有多少节点为 Ready 状态，并且去除包含 NoSchedule 污点的节点。之后将数字写到 `/opt/KUSC00402/kusc00402.txt` <https://edu.51cto.com/lecturer/11062970.html?type=2>
解题：

```
$ kubectl config use-context k8s  
$ kubectl get node | grep -i ready # 记录总数为 A  
$ kubectl describe node | grep Taint | grep NoSchedule # 记录总数为 B  
# 将 A 减 B 的值 x 导入到 /opt/KUSC00402/kusc00402.txt  
$ echo x >> /opt/KUSC00402/kusc00402.txt
```

1.5 节点维护

[← Previous](#)

[☰ Task 2 of 17 ▾](#)

[→ Next](#)

Task weight: 4%

Set configuration context: 

```
[student@node-1] $ | kubectl config use-context  
ek8s
```

Task
Set the node named `ek8s-node-1` as unavailable and reschedule all the pods running on it.

[🚩 Flag this to return to later](#)

[I am satisfied, next →](#)

中文解释：

将 `ek8s-node-1` 节点设置为不可用，然后重新调度该节点上的所有 Pod

解题：

```
$ kubectl config use-context ek8s  
$ kubectl cordon ek8s-node-1  
$ kubectl drain ek8s-node-1 --delete-emptydir-data --ignore-daemonsets --force
```

<https://kubernetes.io/zh/docs/tasks/configure-pod-container/>

1.6 指定节点部署

Task

Schedule a pod as follows:

- Name: `nginx-kusc00401`
- Image: `nginx`
- Node selector: `disk=spinning`

中文解释:

创建一个 Pod，名字为 `nginx-kusc00401`，镜像地址是 `nginx`，调度到具有 `disk=spinning` 标签的节点上，该题可以参考链接：<https://kubernetes.io/zh/docs/concepts/scheduling-eviction/assign-pod-node/>

参考:

<https://kubernetes.io/zh/docs/tasks/configure-pod-container/assign-pods-nodes/>

解题:

```
$ vim pod-ns.yaml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-kusc00401
  labels:
    role: nginx-kusc00401
spec:
  nodeSelector:
    disk: spinning
  containers:
  - name: nginx
    image: nginx
$ kubectl create -f pod-ns.yaml
```

1.7 一个 Pod 多个容器

```
[student@node-1] $ | kubectl config use-context  
k8s
```

Task

Create a pod named `kucc1` with a single app container for each of the following images running inside (there may be between 1 and 4 images specified): `nginx + redis + memcached + consul`.

中文解释：

创建一个 Pod，名字为 `kucc1`，这个 Pod 可能包含 1-4 容器，该题为四个：
`nginx+redis+memcached+consul`

解题：<https://edu.51cto.com/lecturer/11062970.html?type=2>

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: kucc1  
spec:  
  containers:  
    - image: nginx  
      name: nginx  
    - image: redis  
      name: redis  
    - image: memcached  
      name: memcached  
    - image: consul  
      name: consul
```


1.8 Service

Set configuration context:



```
[student@node-1] $ | kubectl config use-context  
k8s
```

Task

Reconfigure the existing deployment `front-end` and add a port specification named `http` exposing port `80/tcp` of the existing container `nginx`.

Create a new service named `front-end-svc` exposing the container port `http`.

Configure the new service to also expose the individual Pods via a NodePort on the nodes on which they are scheduled.

中文解释：

重新配置一个已经存在的 deployment `front-end`，在名字为 `nginx` 的容器里面添加一个端口配置，名字为 `http`，暴露端口号为 `80`，然后创建一个 service，名字为 `front-end-svc`，暴露该 deployment 的 `http` 端口，并且 service 的类型为 `NodePort`。

解题：

本题可以参考：<https://kubernetes.io/docs/concepts/services-networking/connect-applications-service/>

```
$ kubectl edit deploy front-end  
# 添加如下配置，主要是在 name 为 nginx 的容器下
```

```
imagePullPolicy: Always  
name: nginx  
ports:  
- containerPort: 80  
  name: http  
  protocol: TCP
```

添加 service:

```
$
```

```
kubectl expose deploy front-end --name=front-end-svc --port=80 --target-port=http --type=NodePort
```

1.9 Ingress

```
[student@node-1] $ | kubectl config use-context  
t k8s
```

Task

Create a new nginx **Ingress** resource as follows:

- Name: **pong**
- Namespace: **ing-internal**
- Exposing service **hi** on path **/hi** using service port **5678**

The availability of service **hi** can be checked using the following command, which should return **hi** :

```
[student@node-1] $ | curl -kL <INTERNAL_IP>/hi
```

中文解释：

在 ing-internal 命名空间下创建一个 ingress，名字为 pong，代理的 service hi，端口为 5678，配置路径/hi。

验证：访问 curl -kL <INTERNAL_IP>/hi 会返回 hi

解题：

本地可参考：<https://kubernetes.io/zh/docs/concepts/services-networking/ingress/>

```
apiVersion: networking.k8s.io/v1  
kind: Ingress  
metadata:  
  name: pong  
  namespace: ing-internal  
spec:  
  rules:  
  - http:  
    paths:
```

```
- path: /hi
  pathType: Prefix
  backend:
    service:
      name: hi
      port:
        number: 5678
```

1.10 Sidecar

Context

Without changing its existing containers, an existing Pod needs to be integrated into Kubernetes's built-in logging architecture (e.g. `kubectl logs`). Adding a streaming sidecar container is a good and common way to accomplish this requirement.

Task

Add a `busybox` sidecar container to the existing Pod `legacy-app`. The new sidecar container has to run the following command:

```
/bin/sh -c tail -n+1 -f /var/log/legacy-app.log
```

Use a volume mount named `logs` to make the file `/var/log/legacy-app.log` available to the sidecar container.

Don't modify the existing container.
Don't modify the path of the log file, both containers must access it at `/var/log/legacy-app.log`.

中文解释：

添加一个名为 `busybox` 且镜像为 `busybox` 的 sidecar 到一个已经存在的名为 `legacy-app` 的 Pod 上，这个 sidecar 的启动命令为 `/bin/sh, -c, 'tail -n+1 -f /var/log/legacy-app.log'`。

并且这个 sidecar 和原有的镜像挂载一个名为 `logs` 的 volume，挂载的目录为 `/var/log/`
解题：

本题答案: <https://kubernetes.io/zh/docs/concepts/cluster-administration/logging/>

首先将 legacy-app 的 Pod 的 yaml 导出, 大致如下:

```
$ kubectl get po legacy-app -oyaml > c-sidecar.yaml
apiVersion: v1
kind: Pod
metadata:
  name: legacy-app
spec:
  containers:
  - name: count
    image: busybox
    args:
    - /bin/sh
    - -c
    - >
      i=0;
      while true;
      do
        echo "$(date) INFO $i" >> /var/log/legacy-ap.log;
        i=$((i+1));
        sleep 1;
      done
```

再此 yaml 中添加 sidecar 和 volume

```
$ vim c-sidecar.yaml
apiVersion: v1
kind: Pod
metadata:
  name: legacy-app
spec:
  containers:
  - name: count
    image: busybox
    args:
    - /bin/sh
    - -c
    - >
      i=0;
      while true;
      do
        echo "$(date) INFO $i" >> /var/log/legacy-ap.log;
        i=$((i+1));
        sleep 1;
      done
  volumeMounts:
  - name: logs
    mountPath: /var/log
  - name: busybox
    image: busybox
    args: [/bin/sh, -c, 'tail -n+1 -f /var/log/legacy-ap.log']
    volumeMounts:
    - name: logs
      mountPath: /var/log
  volumes:
  - name: logs
    emptyDir: {}
$ kubectl delete -f c-sidecar.yaml ; kubectl create -f c-sidecar.yaml
```

1.11 RBAC

Context

You have been asked to create a new `ClusterRole` for a deployment pipeline and bind it to a specific `ServiceAccount` scoped to a specific namespace.

Task

Create a new `ClusterRole` named `deployment-clusterrole`, which only allows to create the following resource types:

- `Deployment`
- `StatefulSet`
- `DaemonSet`

Create a new `ServiceAccount` named `cicd-token` in the existing namespace `app-team1`.

Bind the new `ClusterRole` `deployment-clusterrole` to the new `ServiceAccount` `cicd-token`, limited to the namespace `app-team1`.

中文解释:

创建一个名为 `deployment-clusterrole` 的 `clusterrole`，该 `clusterrole` 只允许创建 `Deployment`、`Daemonset`、`Statefulset` 的 `create` 操作

在名字为 `app-team1` 的 namespace 下创建一个名为 `cicd-token` 的 `serviceAccount`，并且将上一步创建 `clusterrole` 的权限绑定到该 `serviceAccount`

解题:

可参考: <https://kubernetes.io/zh/docs/reference/access-authn-authz/rbac/>

```
创建 clusterrole
[root@k8s-master01 ~]# cat dp-clusterrole.yaml
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  # "namespace" omitted since ClusterRoles are not namespaced
  name: deployment-clusterrole
rules:
- apiGroups: ["extensions", "apps"]
  #
  # at the HTTP level, the name of the resource for accessing Secret
  # objects is "secrets"
  resources: ["deployments", "statefulsets", "daemonsets"]
```

```
verbs: ["create"]
[root@k8s-master01 ~]# kubectl create -f dp-clusterrole.yaml
clusterrole.rbac.authorization.k8s.io/deployment-clusterrole created
创建 serviceAccount
# kubectl create sa cicd-token -n app-team1
serviceaccount/cicd-token created
```

绑定权限（推荐，节省时间）

```
[root@k8s-master01 ~]# kubectl create rolebinding deployment-rolebinding
--clusterrole=deployment-clusterrole --serviceaccount=app-team1:cicd-token -
n app-team1
```

或者

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: deployment-rolebinding
  namespace: app-team1
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: deployment-clusterrole
subjects:
- kind: ServiceAccount
  name: cicd-token
  namespace: app-team1
```

1.12 <https://www.bilibili.com/lecturer/11062970.html?type=2>

Set configuration context:

```
[student@node-1] $ | kubectl config use-context
hk8s
```

Task

Create a new **NetworkPolicy** named **allow-port-from-namespace** that allows Pods in the existing namespace **internal** to connect to port **9000** of other Pods in the same namespace.

Ensure that the new **NetworkPolicy** :

- does **not** allow access to Pods not listening on port **9000**
- does **not** allow access from Pods not in namespace **internal**

中文解释：

创建一个名字为 **allow-port-from-namespace** 的 **NetworkPolicy**，这个 **NetworkPolicy** 允许 **internal** 命名空间下的 Pod 访问该命名空间下的 **9000** 端口。

并且不允许不是 internal 命名空间的 Pod 访问
不允许访问没有监听 9000 端口的 Pod。

解题：

参考：<https://kubernetes.io/zh/docs/concepts/services-networking/network-policies/>

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-port-from-namespace
  namespace: internal
spec:
  ingress:
  - from:
    - podSelector: {}
    ports:
    - port: 9000
      protocol: TCP
  podSelector: {}
  policyTypes:
  - Ingress
```

1.12.1 NetworkPolicy 此题可能存在的变化

上述的题目是只限制在 internal 命名空间下的，该题可能存在更新。更新如下：

在现有的 namespace my-app 中创建一个名为 allow-port-from-namespace 的 NetworkPolicy 确保这个 NetworkPolicy 允许 namespace my-app 中的 pods 可以连接到 namespace big-corp 中的 8080。

并且不允许不是 my-app 命名空间的 Pod 访问
不允许访问没有监听 8080 端口的 Pod。

所以可以拿着上述的答案，进行稍加修改（注意 namespaceSelector 的 labels 配置，首先需要查看 big-corp 命名空间有没有标签：kubect! get ns big-corp --show-labels 如果有，可以更改 name: big-corp 为查看到的即可。如果没有需要添加一个 label: kubect! label ns big-corp name=big-corp):

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-port-from-namespace
  namespace: my-app
spec:
  egress:
  - to:
    - namespaceSelector:
        matchLabels:
          name: big-corp
    ports:
    - protocol: TCP
      port: 8080
  ingress:
  - from:
    - podSelector: {}
    ports:
    - port: 8080
      protocol: TCP
```

```
podSelector: {}
policyTypes:
- Ingress
- Egress
```

1.12.2 NetworkPolicy 此题可能存在的变化 2

The screenshot shows a task interface for configuring a NetworkPolicy. At the top, there's a green bar with 'time remaining'. Below it are navigation buttons: '← Previo...', 'Task 5 of 16', and '→ Next'. A status bar indicates '问题权重: 7%'. A red box with a warning icon contains the instruction '设置配置环境:' followed by a terminal snippet: '[student@node-1] \$ | kubectl config use-context hk8s'. Below this, the 'Task' section describes the goal: '在现有的 namespace big-corp 中创建一个名为 allow-port-from-namespace 的新 NetworkPolicy。' and '确保新的 NetworkPolicy 允许 namespace internal 中的 Pods 连接到 namespace big-corp 中的 Pods 的端口 9200。' It then lists requirements: '进一步确保新的 NetworkPolicy :', '不允许对没有在监听端口 9200 的 Pods 的访问', and '不允许非来自 namespace internal 中的 Pods 的访问'. At the bottom are buttons: 'Flag this to return to I...' and 'I am satisfied, next →'.

解题:

参考: <https://kubernetes.io/zh/docs/concepts/services-networking/network-policies/>

此题和上题比较比较简单, 只需要允许 internal 命名空间即可

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-port-from-namespace
  namespace: big-corp
spec:
  ingress:
    - from:
        - namespaceSelector:
            matchLabels:
              kubernetes.io/metadata.name: internal
      ports:
        - port: 80
          protocol: TCP
  podSelector: {}
  policyTypes:
    - Ingress
```


1.13 PersistentVolume

Set configuration context:



```
[student@node-1] $ | kubectl config use-context  
hk8s
```

Task

Create a persistent volume with name `app-config`, of capacity `2Gi` and access mode `ReadWriteMany`. The type of volume is `hostPath` and its location is `/srv/app-config`.

中文解释：
<https://edu.51cto.com/lecturer/11062970.html?type=2>
创建一个 pv，名字为 `app-config`，大小为 `2Gi`，访问权限为 `ReadWriteMany`。Volume 的类型为 `hostPath`，路径为 `/srv/app-config`

解题：

参 考： <https://kubernetes.io/docs/tasks/configure-pod-container/configure-persistent-volume-storage/>

```
apiVersion: v1  
kind: PersistentVolume  
metadata:  
  name: app-config  
  labels:  
    type: local  
spec:  
  storageClassName: manual  
  capacity:  
    storage: 2Gi  
  accessModes:  
    - ReadWriteMany  
  hostPath:  
    path: "/srv/app-config"
```

1.14 CSI & PersistentVolumeClaim

```
[student@node-1] $ | kubectl config use-context  
t ok8s
```

Task

Create a new `PersistentVolumeClaim` :

- Name: `pv-volume`
- Class: `csi-hostpath-sc`
- Capacity: `10Mi`

Create a new Pod which mounts the `PersistentVolumeClaim` as a volume:

- Name: `web-server`
- Image: `nginx`
- Mount path: `/usr/share/nginx/html`

Configure the new Pod to have `ReadWriteOnce` access on the volume.

Finally, using `kubectl edit` or `kubectl patch` expand the `PersistentVolumeClaim` to a capacity of `70Mi` and record that change.

中文文档:

创建一个名字为 `pv-volume` 的 `pvc`, 指定 `storageClass` 为 `csi-hostpath-sc`, 大小为 `10Mi` 然后创建一个 Pod, 名字为 `web-server`, 镜像为 `nginx`, 并且挂载该 `PVC` 至 `/usr/share/nginx/html`, 挂载的权限为 `ReadWriteOnce`。之后通过 `kubectl edit` 或者 `kubectl patch` 将 `pvc` 改成 `70Mi`, 并且记录修改记录。

解题:

参 考 : <https://kubernetes.io/docs/tasks/configure-pod-container/configure-persistent-volume-storage/>

创建 `PVC`:

```
apiVersion: v1  
kind: PersistentVolumeClaim  
metadata:  
  name: pv-volume
```

```
spec:
  accessModes:
  - ReadWriteOnce
  resources:
    requests:
      storage: 10Mi
  storageClassName: csi-hostpath-sc
```

创建 Pod:

```
apiVersion: v1
kind: Pod
metadata:
  name: web-server
spec:
  containers:
  - name: nginx
    image: nginx
    volumeMounts:
    - mountPath: "/usr/share/nginx/html"
      name: pv-volume
  volumes:
  - name: pv-volume
    persistentVolumeClaim:
      claimName: pv-volume
```

扩容:

方式一 Patch 命令:

```
kubect1 patch pvc pv-volume -p
'{"spec":{"resources":{"requests":{"storage": "70Mi"}}}}' --record
```

方式二 edit:

```
kubect1 edit pvc pv-volume
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  annotations:
    kubernetes.io/change-cause: 'resize'
    pv.kubernetes.io/bind-completed: "yes"
    pv.kubernetes.io/bound-by-controller: "yes"
    volume.beta.kubernetes.io/storage-provisioner: hostpath.csi.k8s.io
  finalizers:
  - kubernetes.io/pvc-protection
  name: pv-volume
  namespace: default
  resourceVersion: "12950116"
  uid: 1f259dcc-129f-4170-a06a-48fd588b7f5e
spec:
  accessModes:
  - ReadWriteOnce
  resources:
    requests:
      storage: 70Mi
  storageClassName: csi-hostpath-sc
```

1.15 Etcd 备份恢复

No configuration context change required for this item.

Task

First, create a snapshot of the existing `etcd` instance running at `https://127.0.0.1:2379`, saving the snapshot to `/srv/data/etcd-snapshot.db`.

Creating a snapshot of the given instance is expected to complete in seconds.

If the operation seems to hang, something's likely wrong with your command. Use `CTRL + C` to cancel the operation and try again.

Next, restore an existing, previous snapshot located at `/var/lib/backup/etcd-snapshot-previous.db`.

The following TLS certificates/key are supplied for connecting to the server with `etcdctl`:

- CA certificate: `/opt/KUIN00601/ca.crt`
- Client certificate: `/opt/KUIN00601/etcd-client.crt`
- Client key: `/opt/KUIN00601/etcd-client.key`

中文解释：

针对 `etcd` 实例 `https://127.0.0.1:2379` 创建一个快照，保存到 `/srv/data/etcd-snapshot.db`。
在创建快照的过程中，如果卡住了，就键入 `ctrl+c` 终止，然后重试。

然后恢复一个已经存在的快照：`/var/lib/backup/etcd-snapshot-previous.db`

执行 `etcdctl` 命令的证书存放在：

ca 证书：`/opt/KUIN00601/ca.crt`

客户端证书：`/opt/KUIN00601/etcd-client.crt`

客户端密钥：`/opt/KUIN00601/etcd-client.key`

解题:

可参考: <https://kubernetes.io/zh/docs/tasks/administer-cluster/configure-upgrade-etcd/>

```
$ export ETCDCTL_API=3
$ etcdctl --endpoints="https://127.0.0.1:2379" --cacert=/opt/KUIN000601/ca.crt --cert=/opt/KUIN000601/etcd-client.crt --key=/opt/KUIN000601/etcd-client.key snapshot save /srv/data/etcd-snapshot.db
```

还原

```
$ mkdir /opt/backup/ -p
$ cd /etc/kubernetes/manifests ; mv kube-* /opt/backup
$ export ETCDCTL_API=3 etcdctl --endpoints="https://127.0.0.1:2379" --cacert=/opt/KUIN000601/ca.crt --cert=/opt/KUIN000601/etcd-client.crt --key=/opt/KUIN000601/etcd-client.key snapshot restore /var/lib/backup/etcd-snapshot-previous.db --data-dir=/var/lib/etcd-restore
$ vim etcd.yaml
```

将 volume 配置的 path: /var/lib/etcd 改成/var/lib/etcd-restore

```
volumes:
- hostPath:
    path: /etc/kubernetes/pki/etcd
    type: DirectoryOrCreate
  name: etcd-certs
- hostPath:
    path: /var/lib/etcd-restore
```

还原 k8s 组件

```
$ mv /opt/backup/* /etc/kubernetes/manifests
$ systemctl restart kubelet
```

<https://edu.51cto.com/lecturer/11062970.html?type=2>

注意

如果是二进制安装的 etcd, 考试环境的 etcd 可能并非 root 用户启动的, 所以可以先切换到 root 用户 (sudo su -), 然后使用 ps aux | grep etcd 查看启动用户是谁和启动的配置文件是谁 config-file 字段指定, **假设用户是 etcd**。所以如果是二进制安装的 etcd, 执行恢复时需要 root 权限, 所以在恢复数据时, 可以使用 root 用户恢复, 之后更改恢复目录的权限: sudo chown -R etcd:etcd /var/lib/etcd-restore, 然后通过 systemctl status etcd (或者 ps aux | grep etcd) 找到它的配置文件 (**如果没有配置文件, 就可以直接在 etcd 的 service 【通过 systemctl status etcd 即可看到】文件中找到 data-dir 的配置**), 然后更改 data-dir 配置后, 执行 systemctl daemon-reload, 最后使用 etcd 用户 systemctl restart etcd 即可。

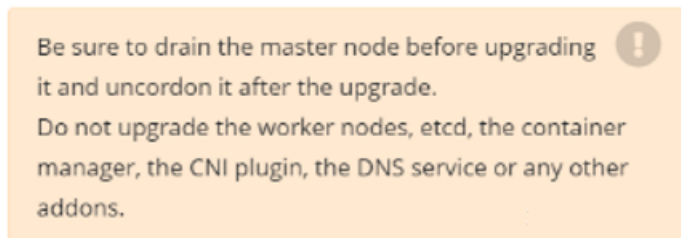
1.16 K8s 升级



Task

Given an existing Kubernetes cluster running version **1.18.8**, upgrade all of the Kubernetes control plane and node components **on the master node only** to version **1.19.0**.

You are also expected to upgrade **kubelet** and **kubectl** on the master node.



<https://edu.51cto.com/lecturer/11062970.html?type=2>
解题:

参考: <https://kubernetes.io/zh/docs/tasks/administer-cluster/kubeadm/kubeadm-upgrade/>

首先腾空节点:

```
# 设置为维护状态
$ kubectl cordon k8s-master
# 驱逐 Pod
$ kubectl drain k8s-master --delete-emptydir-data --ignore-daemonsets --
force
# 之后需要按照题目提示 ssh 到一个 master 节点
$ apt update
$ apt-cache policy kubeadm | grep 1.19.0 # (注意版本的差异, 有可能并非 1.18.8
升级到 1.19)
$ apt-get install kubeadm=1.19.0-00
```

验证升级计划

```
$ kubeadm upgrade plan
```

看到如下信息, 可升级到指定版本

You can now apply the upgrade by executing the following command:

```
kubeadm upgrade apply v1.19.0
```

开始升级 Master 节点

```
$ kubeadm upgrade apply v1.19.0 --etcd-upgrade=false
```

```
[upgrade/successful] SUCCESS! Your cluster was upgraded to "v1.19.0".
```

Enjoy!

```
[upgrade/kubelet] Now that your control plane is upgraded, please
```

proceed with upgrading your kubelets if you haven't already done so.

升级 kubectl 和 kubelet

```
$ apt-get install -y kubelet=1.19.0-00 kubectl=1.19.0-00
$ systemctl daemon-reload
$ systemctl restart kubelet
$ kubectl uncordon k8s-master
node/k8s-master uncordoned
$ kubectl get node
NAME                STATUS    ROLES    AGE   VERSION
k8s-master01        NotReady control-plane,master 11d   v1.19.0
k8s-node01           Ready     <none>    8d    v1.18.8
k8s-node02           Ready     <none>    11d   v1.18.8
$ kubectl get node
NAME                STATUS    ROLES    AGE   VERSION
k8s-master01        Ready     control-plane,master 11d   v1.19.0
k8s-node01           Ready     <none>    8d    v1.18.8
k8s-node02           Ready     <none>    11d   v1.18.8
```

1.16.1 升级到 1.21.1

解题:

参考: <https://kubernetes.io/zh/docs/tasks/administer-cluster/kubeadm/kubeadm-upgrade/>

首先腾空节点:

```
# 设置为维护状态
$ kubectl cordon k8s-master
# 驱逐 Pod
$ kubectl drain k8s-master --delete-emptydir-data --ignore-daemonsets --force
# 之后需要按照题目提示 ssh 到一个 master 节点
$ apt update
$ apt-cache policy kubeadm | grep 1.21.1 # (注意版本的差异, 有可能并非 1.20.1 升级到 1.21.1)
$ apt-get install kubeadm=1.21.1-00
# 验证升级计划
$ kubeadm upgrade plan
# 看到如下信息, 可升级到指定版本
```

开始升级 Master 节点, 注意看题需不需要升级 etcd

```
$ kubeadm upgrade apply v1.21.1 --etcd-upgrade=false -f
```

```
[upgrade/successful] SUCCESS! Your cluster was upgraded to "v1.21.1". Enjoy!
```

```
[upgrade/kubelet] Now that your control plane is upgraded, please proceed with upgrading your kubelets if you haven't already done so.
```

注意

自己的环境升级, 可能会报找不到 coredns 的镜像, 可以使用如下方法解决:

所有节点 `docker pull coredns/coredns:1.8.0`; `docker tag coredns/coredns:1.8.0 registry.cn-hangzhou.aliyuncs.com/google_containers/coredns/coredns:v1.8.0` 然后继续就行。1.8.0 改成你自己 CoreDNS 报错的版本

升级 kubectl 和 kubelet

```
$ apt-get install -y kubelet=1.21.1-00 kubectl=1.21.1-00
$ systemctl daemon-reload
$ systemctl restart kubelet
$ kubectl uncordon k8s-master
node/k8s-master uncordoned
$ kubectl get node
NAME                STATUS    ROLES                  AGE   VERSION
k8s-master01        NotReady control-plane,master   11d   v1.21.1
k8s-node01           Ready     <none>                 8d    v1.12.1
$ kubectl get node
NAME                STATUS    ROLES                  AGE   VERSION
k8s-master01        Ready     control-plane,master   11d   v1.21.1
k8s-node01           Ready     <none>                 8d    v1.20.1
k8s-node02           Ready     <none>                 11d   v1.20.1
```

1.17 集群故障排查 – kubelet 故障



32970.html?type=2

Task

A Kubernetes worker node, named `wk8s-node-0` is in state `NotReady`.

Investigate why this is the case, and perform any appropriate steps to bring the node to a `Ready` state, ensuring that any changes are made permanent.

中文解释：

一个名为 wk8s-node-0 的节点状态为 NotReady，让其恢复至正常状态，并确认所有的更改开机自动完成

解题：

```
$ ssh wk8s-node-0
$ sudo -i
# systemctl status kubelet
# systemctl start kubelet
# systemctl enable kubelet
```

1.18 集群故障排查 – 主节点故障

<https://kubernetes.io/zh/docs/tasks/configure-pod-container/static-pod/>